

Chapter 1: Breast Cancer Detection

```
In [6]: # Let explore the dataset and do a few visualizations  
print(df.loc[10])
```

```
# Print the shape of the dataset  
print(df.shape)
```

```
clump_thickness      1  
uniform_cell_size    1  
uniform_cell_shape    1  
marginal_adhesion    1  
single_epithelial_size 1  
bare_nuclei          1  
bland_chromatin       3  
normal_nucleoli       1  
mitoses              1  
class                2  
Name: 10, dtype: object  
(699, 10)
```

```
In [10]: # Define models to train
models = []
models.append(('KNN', KNeighborsClassifier(n_neighbors = 5)))
models.append(('SVM', SVC()))

# evaluate each model in turn
results = []
names = []

for name, model in models:
    kfold = model_selection.KFold(n_splits=10, random_state = seed)
    cv_results = model_selection.cross_val_score(model, X_train, y_train, cv=kfold)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)
```

KNN: 0.962468 (0.018609)
SVM: 0.958929 (0.029934)

Anaconda Prompt

```
(base) C:\Users\test>D:

(base) D:\>cd D:\Tutorial

(base) D:\Tutorial>conda install numpy
```

```
Anaconda Prompt - jupyter notebook

(base) C:\Users\test>D:

(base) D:\>cd D:\Tutorial

(base) D:\Tutorial>jupyter notebook
[I 18:37:09.670 NotebookApp] The port 8888 is already in use, trying another port.
[I 18:37:10.100 NotebookApp] Loading IPython parallel extension
[I 18:37:10.330 NotebookApp] JupyterLab extension loaded from C:\Users\test\Anaconda3\lib\site-packages\jupyterlab
[I 18:37:10.330 NotebookApp] JupyterLab application directory is C:\Users\test\Anaconda3\share\jupyter\lab
[I 18:37:10.360 NotebookApp] Serving notebooks from local directory: D:\Tutorial
[I 18:37:10.360 NotebookApp] The Jupyter Notebook is running at:
[I 18:37:10.360 NotebookApp] http://localhost:8889/?token=b9b0fa73e9ce7368ed96458fa0e133994e987c4befb387dc
[I 18:37:10.360 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[C 18:37:10.380 NotebookApp]

Copy/paste this URL into your browser when you connect for the first time,
to login with a token:
    http://localhost:8889/?token=b9b0fa73e9ce7368ed96458fa0e133994e987c4befb387dc
[I 18:37:10.660 NotebookApp] Accepting one-time-token-authenticated connection from ::1
```

```
In [1]: import sys
import scipy
import numpy
import matplotlib
import pandas
import sklearn

print('Python: {}'.format(sys.version))
print('scipy: {}'.format(scipy.__version__))
print('numpy: {}'.format(numpy.__version__))
print('matplotlib: {}'.format(matplotlib.__version__))
print('pandas: {}'.format(pandas.__version__))
print('sklearn: {}'.format(sklearn.__version__))

Python: 3.6.6 |Anaconda, Inc.| (default, Jun 28 2018, 11:27:44) [MSC v.1900 64
bit (AMD64)]
scipy: 1.1.0
numpy: 1.15.0
matplotlib: 2.2.2
pandas: 0.23.3
sklearn: 0.19.1
```

```
In [2]: import numpy as np
        from sklearn import preprocessing, cross_validation
        from sklearn.neighbors import KNeighborsClassifier
        from sklearn.svm import SVC
        from sklearn import model_selection
        from sklearn.metrics import classification_report, accuracy_score
        from pandas.plotting import scatter_matrix
        import matplotlib.pyplot as plt
        import pandas as pd
```

C:\Users\test\Anaconda3\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
"This module will be removed in 0.20.", DeprecationWarning)

```
In [3]: # Load Dataset
        url = "https://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/breast-cancer-wisconsin.data"
        names = ['id', 'clump_thickness', 'uniform_cell_size', 'uniform_cell_shape',
                 'marginal_adhesion', 'single_epithelial_size', 'bare_nuclei',
                 'bland_chromatin', 'normal_nucleoli', 'mitoses', 'class']
        df = pd.read_csv(url, names=names)
```

```
➤ In [4]: # Preprocess the data
        df.replace('?', -99999, inplace=True)
        print(df.axes)

        df.drop(['id'], 1, inplace=True)

        # Print the shape of the dataset
        print(df.shape)

        [RangeIndex(start=0, stop=699, step=1), Index(['id', 'clump_thickness', 'uniform_cell_size', 'uniform_cell_shape',
        'marginal_adhesion', 'single_epithelial_size', 'bare_nuclei',
        'bland_chromatin', 'normal_nucleoli', 'mitoses', 'class'],
        dtype='object')]
        (699, 10)
```

```
In [5]: # Do dataset visualizations
print(df.loc[6])
```

```
clump_thickness      1
uniform_cell_size    1
uniform_cell_shape   1
marginal_adhesion    1
single_epithelial_size 2
bare_nuclei          10
bland_chromatin       3
normal_nucleoli       1
mitoses              1
class                2
Name: 6, dtype: object
```

In [6]: # Do dataset visualizations

```
print(df.loc[6])
```

```
print(df.describe())
```


```
clump_thickness      1
uniform_cell_size    1
uniform_cell_shape    1
marginal_adhesion     1
single_epithelial_size 2
bare_nuclei          10
bland_chromatin       3
normal_nucleoli       1
mitoses              1
class                 2
Name: 6, dtype: object
```

	clump_thickness	uniform_cell_size	uniform_cell_shape \
count	699.000000	699.000000	699.000000
mean	4.417740	3.134478	3.207439
std	2.815741	3.051459	2.971913
min	1.000000	1.000000	1.000000
25%	2.000000	1.000000	1.000000
50%	4.000000	1.000000	1.000000
75%	6.000000	5.000000	5.000000
max	10.000000	10.000000	10.000000


	marginal_adhesion	single_epithelial_size	bland_chromatin \
count	699.000000	699.000000	699.000000
mean	2.806867	3.216023	3.437768
std	2.855379	2.214300	2.438364
min	1.000000	1.000000	1.000000
25%	1.000000	2.000000	2.000000
50%	1.000000	2.000000	3.000000
75%	4.000000	4.000000	5.000000
max	10.000000	10.000000	10.000000

	normal_nucleoli	mitoses	class
count	699.000000	699.000000	699.000000
mean	2.866953	1.589413	2.689557
std	3.053634	1.715078	0.951273
min	1.000000	1.000000	2.000000
25%	1.000000	1.000000	2.000000
50%	1.000000	1.000000	2.000000
75%	4.000000	1.000000	4.000000
max	10.000000	10.000000	4.000000

```
In [7]: # Plot histograms for each variable
df.hist(figsize = (10, 10))
plt.show()
```



```
In [8]: # Create scatter plot matrix
scatter_matrix(df, figsize = (18,18))
plt.show()
```



```
In [9]: # Create X and Y datasets for training
X = np.array(df.drop(['class'], 1))
y = np.array(df['class'])

X_train, X_test, y_train, y_test = cross_validation.train_test_split(X, y, test_size=0.2)
```



```
In [10]: # Testing Options
        seed = 8
        scoring = 'accuracy'
```

```
In [11]: # Define models to train
models = []
models.append(('KNN', KNeighborsClassifier(n_neighbors = 5)))
models.append(('SVM', SVC()))

# evaluate each model in turn
results = []
names = []

for name, model in models:
    kfold = model_selection.KFold(n_splits=10, random_state = seed)
    cv_results = model_selection.cross_val_score(model, X_train, y_train, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)

KNN: 0.966039 (0.018616)
SVM: 0.955292 (0.021477)
```

```
In [11]: # Define models to train
models = []
models.append(('KNN', KNeighborsClassifier(n_neighbors = 5)))
models.append(('SVM', SVC()))

# evaluate each model in turn
results = []
names = []

for name, model in models:
    kfold = model_selection.KFold(n_splits=10, random_state = seed)
    cv_results = model_selection.cross_val_score(model, X_train, y_train, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    msg = "%s: %f (%f)" % (name, cv_results.mean(), cv_results.std())
    print(msg)

KNN: 0.966039 (0.018616)
SVM: 0.955292 (0.021477)
```

```
In [11]: # Make predictions on validation dataset
```

```
for name, model in models:  
    model.fit(X_train, y_train)  
    predictions = model.predict(X_test)  
    print(name)  
    print(accuracy_score(y_test, predictions))  
    print(classification_report(y_test, predictions))
```

KNN

0.9785714285714285

	precision	recall	f1-score	support
2	0.98	0.99	0.98	95
4	0.98	0.96	0.97	45
avg / total	0.98	0.98	0.98	140

SVM

0.9571428571428572

	precision	recall	f1-score	support
2	1.00	0.94	0.97	95
4	0.88	1.00	0.94	45
avg / total	0.96	0.96	0.96	140

```
In [11]: # Make predictions on validation dataset
```

```
for name, model in models:  
    model.fit(X_train, y_train)  
    predictions = model.predict(X_test)  
    print(name)  
    print(accuracy_score(y_test, predictions))  
    print(classification_report(y_test, predictions))
```

KNN

0.9785714285714285

	precision	recall	f1-score	support
2	0.98	0.99	0.98	95
4	0.98	0.96	0.97	45
avg / total	0.98	0.98	0.98	140

SVM

0.9571428571428572

	precision	recall	f1-score	support
2	1.00	0.94	0.97	95
4	0.88	1.00	0.94	45
avg / total	0.96	0.96	0.96	140

```
In [13]: clf = SVC()
```

```
clf.fit(X_train, y_train)  
accuracy = clf.score(X_test, y_test)  
print(accuracy)  
  
example_measures = np.array([[4,2,1,1,1,2,3,2,1]])  
example_measures = example_measures.reshape(len(example_measures), -1)  
prediction = clf.predict(example_measures)  
print(prediction)
```

0.95

[2]

```
In [12]: clf = SVC()

clf.fit(X_train, y_train)
accuracy = clf.score(X_test, y_test)
print(accuracy)

example_measures = np.array([[4,2,1,1,1,2,3,2,1]])
example_measures = example_measures.reshape(len(example_measures), -1)
prediction = clf.predict(example_measures)
print(prediction)
```

```
In [13]: clf = SVC()

clf.fit(X_train, y_train)
accuracy = clf.score(X_test, y_test)
print(accuracy)

example_measures = np.array([[4,2,1,1,1,2,3,2,1]])
example_measures = example_measures.reshape(len(example_measures), -1)
prediction = clf.predict(example_measures)
print(prediction)
```

```
0.95
[2]
```